

Appendix: Robustness and Falsification Checks

Chaitanya Venkateswaran

25th January 2026

A HAC Bandwidth Robustness

Baseline inference uses Newey–West heteroskedasticity- and autocorrelation-consistent (HAC) standard errors with a horizon-dependent bandwidth $\ell_h = \max(1, h)$. Figure 1 compares the baseline impulse response to estimates obtained using fixed HAC bandwidths $\ell \in \{4, 8, 12\}$ and a data-driven automatic bandwidth rule. The impulse response profiles are nearly identical across bandwidth choices, indicating that the qualitative pattern and magnitudes are not driven by the particular HAC bandwidth specification; bandwidth choice primarily affects the width of confidence intervals rather than the estimated dynamics.

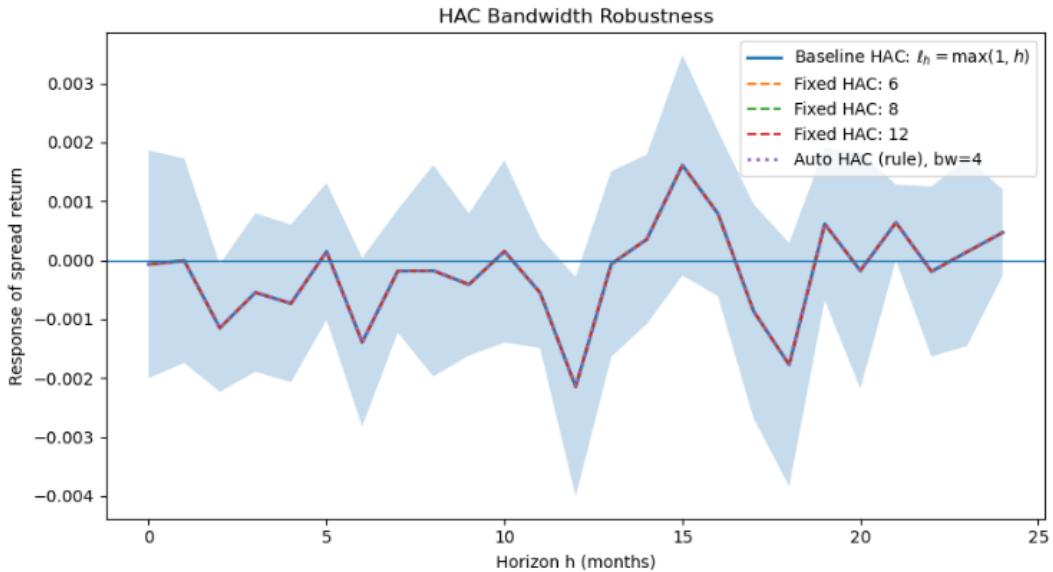


Figure 1: HAC Bandwidth Robustness: Baseline $\ell_h = \max(1, h)$ vs. Automatic + Fixed $\ell \in \{4, 8, 12\}$

B Placebo / Falsification Test

To assess whether the baseline impulse responses could arise mechanically from the local projection specification, I implement a falsification test by replacing the inflation forecast error with a placebo shock. The placebo is constructed to have no economic content for municipal bond pricing (e.g., a randomly permuted shock or a lead of the shock). Figure 2 shows that the placebo generates no stable or systematic dynamic response pattern, supporting the interpretation that the baseline results reflect economically meaningful comovement between inflation forecast errors and relative green municipal bond performance rather than spurious dynamics.

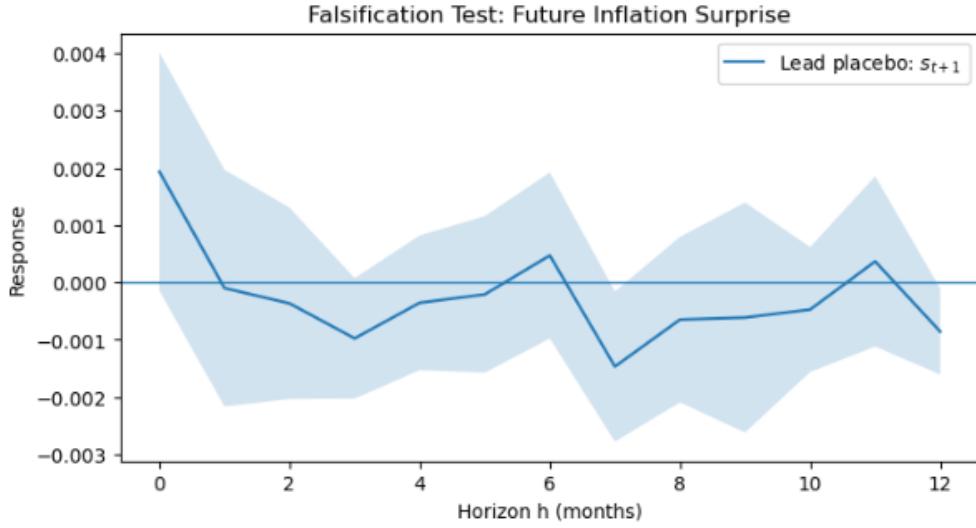


Figure 2: Placebo / Falsification Test: Impulse Responses Using a Placebo Shock

C Lag-Length Robustness

The baseline specification includes $L = 6$ lags of the return spread and the inflation forecast error. To assess sensitivity to this choice, Figure 3 reports impulse responses estimated under alternative lag lengths $L \in \{3, 6, 12\}$. The timing, sign, and overall magnitude of the response are similar across specifications, indicating that the main conclusions are not sensitive to lag length. As expected in a short monthly sample, longer lag specifications can widen confidence intervals at longer horizons, but the qualitative pattern remains stable.




Figure 3: Lag-Length Robustness: Impulse Responses for $L \in \{3, 6, 12\}$